

bayesian_lasso_with_me

December 6, 2023

The goal of this notebook is to implement the Bayesian LASSO method for a 1D problem.

```
[1]: import numpy as np
import matplotlib.pyplot as plt

from scipy.stats import recipinvgauss
import scipy.sparse as sps
from fastprogress import progress_bar


from IPython.display import clear_output, DisplayHandle
def update_patch(self, obj):
    clear_output(wait=True)
    self.display(obj)
DisplayHandle.update = update_patch

from runningstatistics import StatsTracker
import jlinops
```

1 Make toy problem

```
[2]: ground_truth = jlinops.piecewise_constant_1d_test_problem()
n = len(ground_truth)
np.random.seed(0)
noise_stdev = 0.05
noise_var = noise_stdev**2
noisy_signal = ground_truth + noise_stdev*np.random.normal(size=n)
grid = np.arange(n)
```

```
[3]: plt.plot(grid, ground_truth, label="Ground truth", color="C0")
plt.scatter(grid, noisy_signal, marker="x", label="Data", color="C1", alpha=1.
    ↪0, s=0.5)
plt.legend()
plt.show()
```



```
[4]: # Define forward operator and regularization matrix
F = jlinops.MatrixLinearOperator(sps.eye(n))
R, _ = jlinops.first_order_derivative_1d(n, boundary="none")
R = jlinops.MatrixLinearOperator(R)

# Set regularization lambda
reg_lambda = 1e2
```


2 MAP estimate of the L1 problem

By L1 problem, I mean solving

$$\operatorname{argmin}_x \left\{ \frac{1}{2\sigma^2} \|x - y\|_2^2 + \lambda \|Rx\|_1 \right\} = \operatorname{argmin}_x \left\{ \frac{1}{2} \|x - y\|_2^2 + (\lambda\sigma^2) \|Rx\|_1 \right\}.$$

```
[5]: # Solution is given by evaluating the proximal operator of the TV norm. This
      ↪ code uses a FDGP method
fdgp_map_result = jlinops.prox_tv1d_norm(noisy_signal,
      ↪ lam=noise_var*reg_lambda, iterations=1000)
```

```
[6]: plt.scatter(grid, noisy_signal, marker="x", label="Data", color="C1", alpha=1.
      ↪0, s=0.5)
plt.plot(grid, fdgp_map_result, label="MAP estimate of L1 problem", color="C2")
plt.legend()
plt.show()
```




```
[7]: reg_lambda
```

```
[7]: 100.0
```

```
[8]: # Solution is given by evaluating the proximal operator of the TV norm. This
      ↪code uses a FDGP method
fdgp_map_result = jlinops.prox_tv1d_norm(noisy_signal, lam=noise_var*25,
      ↪iterations=100)
```

```
[9]: plt.scatter(grid, noisy_signal, marker="x", label="Data", color="C1", alpha=1.
      ↪0, s=0.5)
plt.plot(grid, fdgp_map_result, label="MAP estimate of L1 problem", color="C2")
plt.legend()
plt.show()
```



3 Sample the posterior of the Gaussian model

$$-\log \pi(x) = \left\{ \frac{1}{2\sigma^2} \|x - y\|_2^2 + \frac{\lambda}{2} \|Rx\|_2^2 \right\} + C$$


```
[10]: def gauss_posterior_summary(F, R, y, noise_var=1.0, reg_lambda=1e1):
    """Computes posterior mean and stdev.
    """
    Q = sps.csc_matrix((1/noise_var)*(F.A.T @ F.A) + reg_lambda*(R.A.T @ R.A))
    Q = jlinops.MatrixLinearOperator(Q)
    Linv = jlinops.BandedCholeskyFactorInvOperator(Q)
    mean = Linv.T @ (Linv @ ((1/noise_var)*F.T @ y) )

    # Get diagonal entries of Qinvs
    Qinvs = Linv.T @ Linv
    var = jlinops.black_box_diagonal(Qinvs)
    stdev = np.sqrt(var)
    return mean, stdev
```

```
[11]: gauss_mu, gauss_sigmas = gauss_posterior_summary(F, R, noisy_signal,
↳noise_var=noise_var, reg_lambda=1e4)
```

<IPython.core.display.HTML object>

```
[12]: #plt.scatter(grid, noisy_signal, marker="x", label="Data", color="C1", alpha=1.
↳0, s=0.5)
plt.plot(grid, gauss_mu, label="Mean vector", color="C2")
plt.fill_between(grid, gauss_mu - 3*gauss_sigmas, gauss_mu + 3*gauss_sigmas,
↳color="C4", alpha=0.5, label="3 sigma band")
plt.legend()
plt.show()
```



4 Bayesian LASSO

```
[13]: class BayesianLASSOGibbsSampler:
    """Implements the Bayesian LASSO hierarchical sampler for the L1 problem.
    """

    def __init__(self, F, R, y, noise_var=1.0):
```

```

self.F = F
self.R = R
self.y = y
self.noise_var = noise_var
self.reg_lambda = reg_lambda

def sample(self, n_samples, x0=None, n_burn=0, theta_tol=1e-2, lam0=None,
↳ lam_update_freq=25):
    """Runs the Gibbs sampler.
    """

    # Initialize
    if x0 is None:
        x = np.zeros(self.F.shape[1])
    else:
        x = x0

    if lam0 is None:
        lam = 1.0
    else:
        lam = lam0

    # Create trackers
    x_tracker = StatsTracker(self.F.shape[1])
    theta_tracker = StatsTracker(self.R.shape[0])

    # For taking care of lambda updates
    lam_update_fn = lambda theta_ss_est: np.sqrt( 2*self.R.shape[0]/
↳ theta_ss_est )
    theta_sum_tracker = StatsTracker((1,))
    lam_hist = [lam]
    theta_sums_all = []

    # Run the sampler
    for j in progress_bar(range(n_samples+n_burn)):

        # Update theta
        theta = self.sample_theta(x, tol=theta_tol)

        # Update x
        x = self.sample_x(theta)

        # For taking care of lambda updates
        theta_sums_all.append(theta.sum())
        theta_sum_tracker.push(theta.sum())

```

```

    if (j < n_burn) and (j%lam_update_freq == 0):

        # Get new lambda
        lam = lam_update_fn(theta_sum_tracker.mean()[0])
        lam_hist.append(lam)

        # Reset theta sum tracker
        theta_sum_tracker = StatsTracker((1,))

    # Push to tracker
    if j >= n_burn:
        x_tracker.push(x)
        theta_tracker.push(theta)

results = {
    "x_tracker": x_tracker,
    "theta_tracker": theta_tracker,
    "lam_hist": np.asarray(lam_hist),
    "theta_sums_all": np.asarray(theta_sums_all),
}

return results

def sample_x(self, theta):
    """Given local variances theta, draws a sample for x.
    """

    Q = (1.0/self.noise_var)*(self.F.A.T @ self.F.A) + (1/2)*(self.R.A.T @
↪( sps.diags(1.0/theta) @ self.R.A ) )

    # # Bad way
    # Qinvs = np.linalg.inv(Q.toarray())
    # mean = Qinvs @ ((1.0/self.noise_var)*self.F.T @ self.y )
    # sample = np.random.multivariate_normal(mean, Qinvs)

    # Good way
    Q = sps.csc_matrix(Q)
    Q = jlinops.MatrixLinearOperator(Q)
    Linv = jlinops.BandedCholeskyFactorInvOperator(Q)
    mean = Linv.T @ (Linv @ ((1.0/self.noise_var)*self.F.T @ self.y ) )
    sample = mean + ( Linv.T @ np.random.normal(size=Q.shape[0]) )

    return sample

```

```

def sample_theta(self, x, tol=1e-2):
    """Given x, draws a sample for the thetas.
    """

    # Get Rx
    Rx = self.R @ x

    # Make output array
    sample = np.zeros(self.R.shape[0])

    # Need to check where Rx is close to zero, so we can sample from
    ↪ exponential there instead
    idx_too_small = np.where(np.abs(Rx) < tol)
    idx_fine = np.where(np.abs(Rx) >= tol)

    # Break into two parts
    Rx_too_small = Rx[idx_too_small]
    Rx_fine = Rx[idx_fine]

    # For the components near zero, sample from the exponential
    theta_from_too_small = np.random.exponential(scale=1.0/self.reg_lambda,
    ↪ size=len(Rx_too_small)))

    # For the components not near zero, sample from the inverse Gaussian
    theta_from_fine = recipinvgauss.rvs(mu=1.0/(self.reg_lambda*np.
    ↪ abs(Rx_fine))), scale=1.0/(self.reg_lambda**2))

    # Put all into one array
    sample[idx_too_small] = theta_from_too_small
    sample[idx_fine] = theta_from_fine

    assert np.all(sample > 0), "some thetas are no positive!"

    return sample

```

```

[14]: lasso_sampler = BayesianLASSOGibbsSampler(F, R, noisy_signal,
    ↪ noise_var=noise_var)

```

```

[15]: sampling_result = lasso_sampler.sample(n_samples=300, n_burn=500,
    ↪ lam0=100*reg_lambda, lam_update_freq=50)

```

<IPython.core.display.HTML object>

```


[16]: sampling_result["lam_hist"]

```

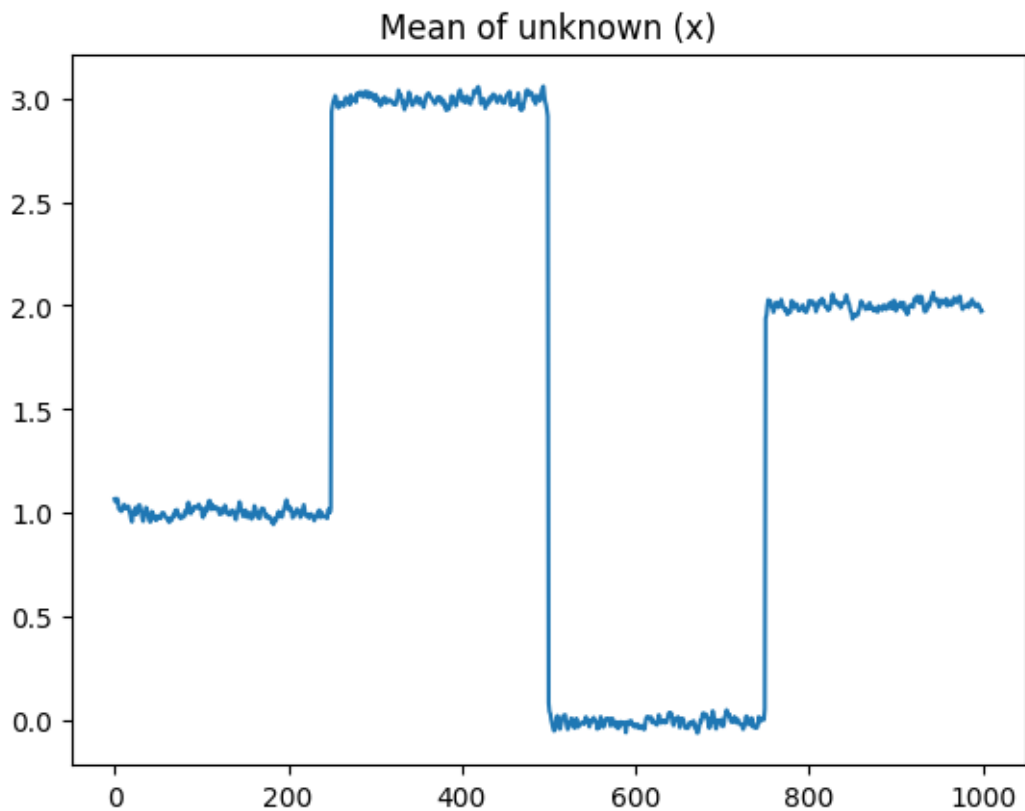


```
[16]: array([10000.      ,  13.81413769,  25.30163257,  25.17612948,  
            25.15326904,  25.16159668,  24.83322155,  25.01487996,  
            25.14007003,  25.01602822,  25.17880348])
```


```
[17]: plt.plot(sampling_result["theta_sums_all"])  
plt.show()
```




```
[18]: plt.plot(sampling_result["x_tracker"].mean())  
plt.title("Mean of unknown (x)")  
plt.show()
```




```
[19]: plt.plot(sampling_result["theta_tracker"].mean())  
plt.title("Mean of local variance parameters (theta)")  
plt.show()
```



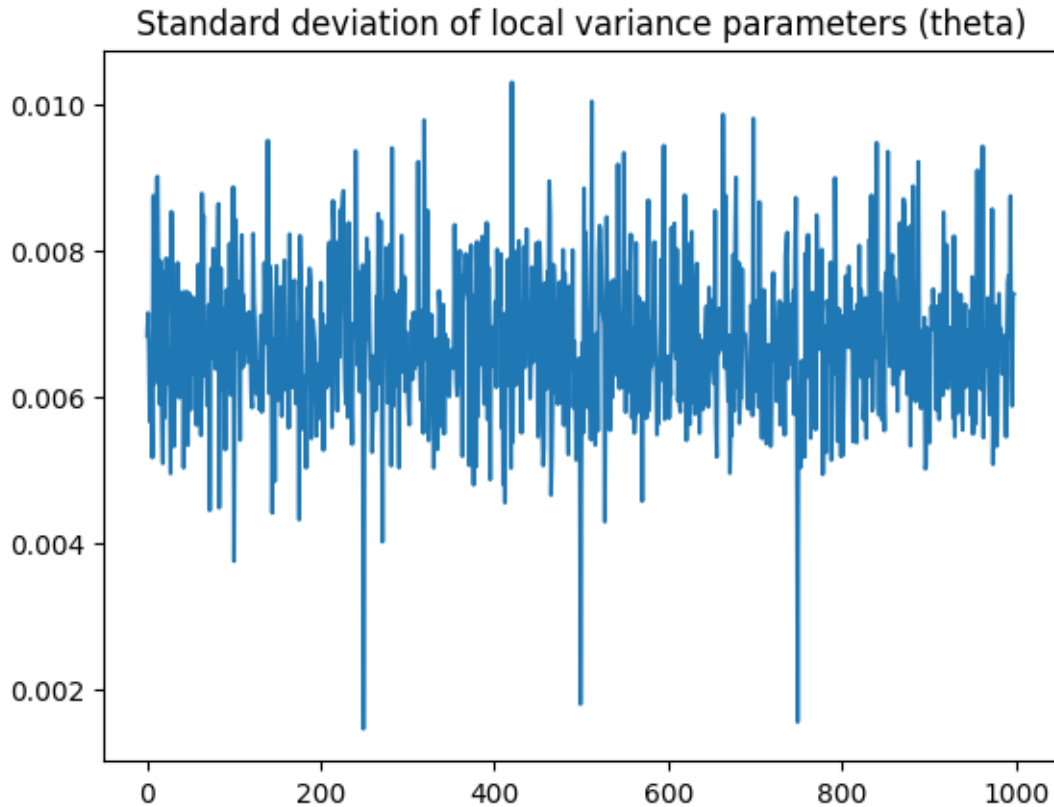
```
[20]: plt.plot(sampling_result["x_tracker"].stdev())  
plt.title("Standard deviation of unknown (x)")  
plt.show()
```



```
[25]: #plt.scatter(grid, noisy_signal, marker="x", label="Data", color="C1", alpha=1.
      ↪0, s=0.5)
mu, stdev = sampling_result["x_tracker"].mean(), sampling_result["x_tracker"].
      ↪stdev()
plt.plot(grid, mu, label="Mean vector", color="C2")
plt.fill_between(grid, mu - 3*stdev, mu + 3*stdev, color="C4", alpha=0.5,
      ↪label="3 sigma band")
plt.legend()
plt.show()
```



```
[26]: plt.plot(sampling_result["theta_tracker"].stdev())  
plt.title("Standard deviation of local variance parameters (theta)")  
plt.show()
```

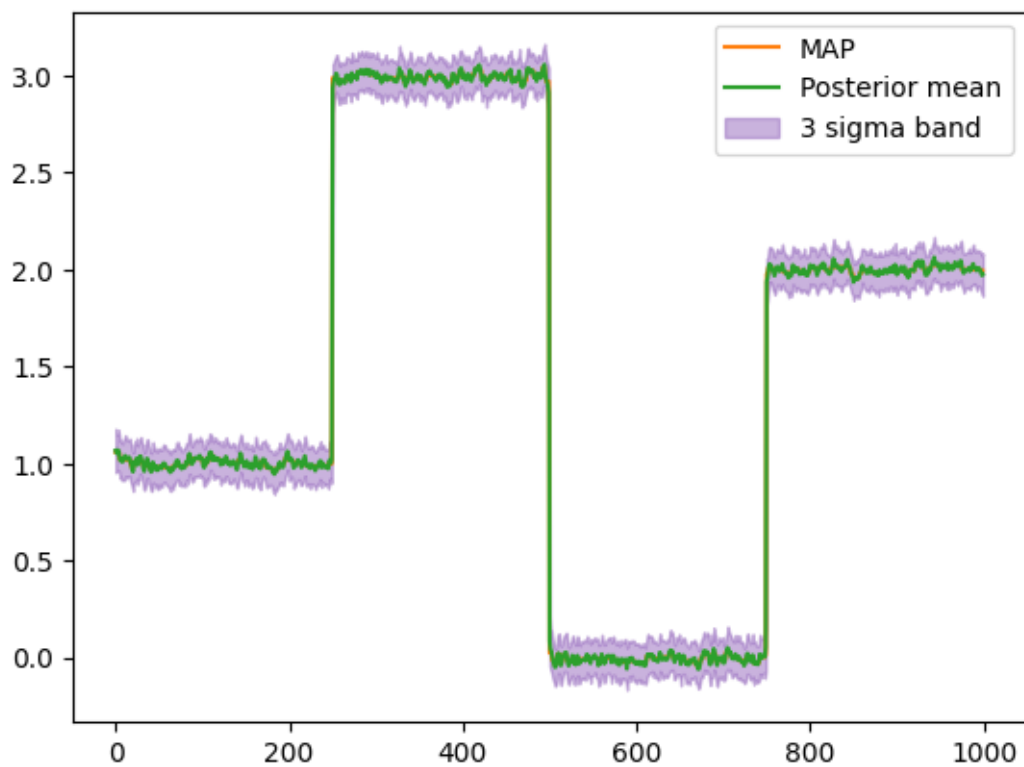


5 Comparison with other methods

```
[27]: # Get MAP point for the lambda chosen during sampling
fdgp_map_result = jlinops.prox_tv1d_norm(noisy_signal,
    ↪ lam=noise_var*sampling_result["lam_hist"][-1], iterations=1000)
#plt.scatter(grid, noisy_signal, marker="x", label="Data", color="C1", alpha=1.
    ↪ 0, s=0.5)


mu, stdev = sampling_result["x_tracker"].mean(), sampling_result["x_tracker"].
    ↪ stdev()

plt.plot(grid, fdgp_map_result, label="MAP", color="C1")
plt.plot(grid, mu, label="Posterior mean", color="C2")
plt.fill_between(grid, mu - 3*stdev, mu + 3*stdev, color="C4", alpha=0.5,
    ↪ label="3 sigma band")
plt.legend()
plt.show()
```



```
[28]: plt.plot( fdgp_map_result - mu )  
      plt.plot()
```

[28]: []



[]:

[]:

[]:

[]:

[]: